

Problem Set 2 for EE227C (Spring 2018): Convex Optimization and Approximation

Instructor: Moritz Hardt

Email: hardt+ee227c@berkeley.edu

Graduate Instructor: Max Simchowitz

Email: msimchow+ee227c@berkeley.edu

March 1, 2018

Problem 1: Backtracking Line Search

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be an m -strongly convex, M -smooth (and thus differentiable) function with global minimum x^* . Consider the following algorithm:

Initialize with an arbitrary $x_0 \in \mathbb{R}^n$, and fix parameters $\alpha \in (0, 1/2)$, $\beta \in (0, 1)$. Then at each step $t = 1, 2, \dots$, do the following:

(a) Let $g_t = \nabla f(x_t)$.

(b) For $k = \{0, 1, \dots\}$ in sequence, check if the following “sufficient decrease” condition holds:

$$f(x - tg_t) \leq f(x) - \alpha\beta^k \cdot \|g_t\|^2 \quad (1)$$

Assuming that this condition holds for some k (you will show this), set $\eta_t = \beta^k$.

(c) Set $x_t \leftarrow x_{t-1} - \eta_t g_t$

(A) Show that condition 1 holds for all $t \in (0, 1/M]$.

(B) Show that $\eta_t \geq \min\{1, \beta/M\}$. Conclude that step (b) of the above algorithm always terminates.

(C) Using part b, show that

$$f(x_t - \eta_t g_t) \leq f(x) - \alpha \min\{1, \frac{\beta}{M}\} \|\nabla f(x_t)\|^2 \quad (2)$$

(D) Show that there is a constant $C = C(\alpha, \beta, M, m) < 1$ such

$$f(x_t - \eta_t g_t) - f(x) \leq C(\alpha, \beta, M, m) \cdot (f(x_t) - f(x_t)) \quad (3)$$

Problem 2: Random Descent Directions

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be an m -strongly convex, M -smooth (and thus differentiable) function with global minimum x^* . Consider the following algorithm: Initialize with an arbitrary $x_0 \in \mathbb{R}^n$. Then at each step $t = 1, 2, \dots$, do the following:

- (a) Choose $g_t \stackrel{\text{unif}}{\sim} \mathcal{S}^{n-1}$ (equivalently, g_t has the distribution of $\frac{g}{\|g\|}$, where $g \sim \mathcal{N}(0, I_n)$).
- (b) Compute a step size $\eta_t := \min_{\eta \geq 0} f(x_{t-1} - \eta g_t)$.
- (c) set $x_t \leftarrow x_{t-1} - \eta g_t$

- (A) Prove that the above algorithm is a (non-strict) descent method; that is $f(x_t)$ is non-increasing in t . Also prove that unless $x_t = x_*$, $f(x_{t+1}) < f(x_t)$ with probability $1/2$.
- (B) Prove that there exists a numerical constant C such that, if

$$t \geq T(\epsilon) := Cn \cdot \frac{M}{m} \log\left(\frac{f(x_0) - f(x^*)}{\epsilon}\right), \quad (4)$$

then $\mathbb{E}[f(x_t) - f(x^*)] \leq \epsilon$.

- (C) Amend the stated algorithm to use line search instead of solving for the exactly-optimal step size. Are the rates qualitatively similar?

Problem 3: Sh*t about Quadratics

In this problem, you are going to test the sharpness of our upper and lower bounds for quadratics on a randomly generated instance. Fix $n = 500$. We define the distribution over PSD matrix $\mathcal{D}(\epsilon)$:

Definition 0.1. $\mathcal{D}(\epsilon)$ is a distribution of matrix $\mathbf{M} = \mathbf{M}^\top$, defined as follows. Let $\mathbf{X} \in \mathbb{R}^{n \times n}$ denote a matrix with i.i.d $\mathcal{N}(0, 1)$ entries. Generate a random vector \mathbf{u} uniformly from the unit sphere. Define the matrix $\mathbf{M} = \frac{1}{\sqrt{2n}}(\mathbf{X} + \mathbf{X}^\top) + (1 + \epsilon)\mathbf{u}\mathbf{u}^\top$.

Now, for each $\epsilon \in \mathcal{S} := \{1, .5, .2, .1, .05\}$, do the following

- (A) Conduct trials $t = 1, 2, \dots, 10$.
 - (A.1) Generate $\mathbf{M} \sim \mathcal{D}(\epsilon)$ as above, and a random vector \mathbf{v} uniformly on the unit sphere.
 - (A.2) Set $\gamma = 2\lambda_{\max}(\mathbf{W}) - \lambda_2(\mathbf{W})$, and define the matrix $\mathbf{N} = \gamma I - \mathbf{M}$. Definally, define the function $\mathbf{f}(x) = \min_x x^\top \mathbf{N} x - 2\langle \mathbf{v}, x \rangle$. What is the condition number of \mathbf{N} ?
 - (A.3) Setting $x_0 = 0$, run gradient descent, a heavy-ball method or nesterov method to solve $\min_x \mathbf{f}(x)$ for a good number of iterations (use your discretion). You may compute the eigenvalues of \mathbf{N} to tune your step parameters.

- (A.4) For both gradient descent and heavy-ball, record for each trial iteration s , the difference between $\mathbf{f}(x_s) - \min_x \mathbf{f}(x)$ for each iteration.
- (A.5) Using the step sizes, largest/smallest eigenvalues of \mathbf{N} , and the initial point $x_0 = 0$, compute a worst case upper bound for $\mathbf{f}(x_s) - \min_x \mathbf{f}(x)$ for each iteration s of gradient descent and the heavy ball method.
- (A.6) Run gradient descent, but this time compute the optimality gap unising “best” iterate in the Krylov space. THat is, compute

$$\min_{x \in \text{span}(x_1, \dots, x_s)} \mathbf{f}(x) - \min_x \mathbf{f}(x) \quad (5)$$

- (A.7) After each trial, you should have a list of 5 values for each iterate s : an upper bound for gradient descent, the rate actually attained by gradient descent, an upper bound for heavy ball/nesterov, the rate actualy attained by heavy ball/nesterov, and the “optimal” krylov algorithm,
- (B) For each of the lists above, average all 10 trials and plot them on the same plot. How sharp are the upper bounds?